

# ASUCA on GPU: Uncompromising Hybrid Port for Physical Core of Japanese Weather Model

**Michel Müller**

Typhoon Computing, Zurich Switzerland  
michel@typhooncomputing.com

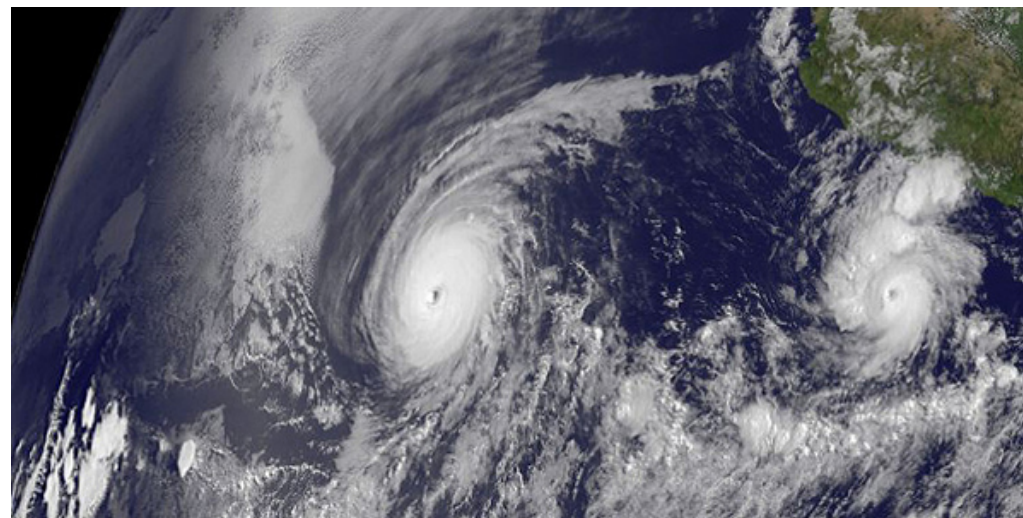
Supervised by

**Dr. Naoya Maruyama**  
RIKEN Advanced Institute for Computational Science

**Dr. Tabito Hara**  
Japan Meteorological Agency

**Dr. Takashi Shimokawabe**  
Tokyo Institute of Technology

**Prof. Dr. Takayuki Aoki**  
Tokyo Institute of Technology



Creative Commons: Nasa Goddard Space Flight Centre, 2010



東京工業大学  
Tokyo Institute of Technology



# ASUCA on GPU: Uncompromising Hybrid Port for Physical Core of Japanese Weather Model

**Michel Müller**

Typhoon Computing, Zurich Switzerland  
michel@typhooncomputing.com

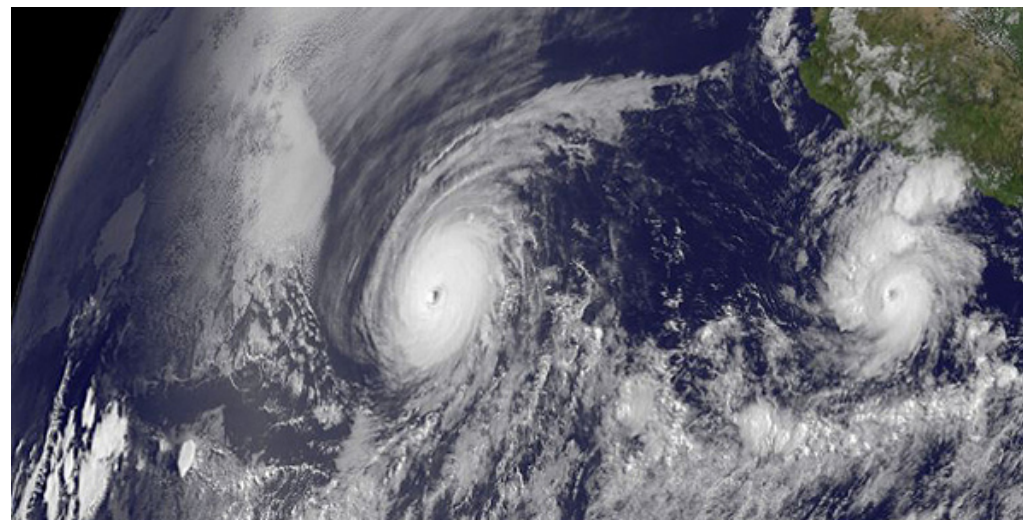
Supervised by

**Dr. Naoya Maruyama**  
RIKEN Advanced Institute for Computational Science

**Dr. Tabito Hara**  
Japan Meteorological Agency

**Dr. Takashi Shimokawabe**  
Tokyo Institute of Technology

**Prof. Dr. Takayuki Aoki**  
Tokyo Institute of Technology



Creative Commons: Nasa Goddard Space Flight Centre, 2010



東京工業大学  
Tokyo Institute of Technology



Motivation



Hybrid Fortran



ASUCA PP Implementation

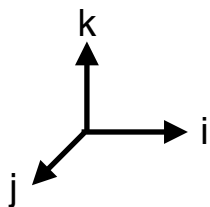


Results & Discussion



Conclusion

- Goals for ASUCA GPGPU portation of physical core:
  - Eliminate host-to-device communication with GPGPU version of the Dynamical Core
  - Gain execution time speedups
  - Portation must remain CPU compatible
  - Portation must remain performant on CPU



# Introduction | Hybrid Fortran

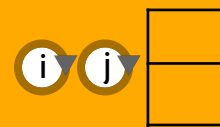
OpenACC  
approach

Original (CPU):  
Parallelization  
at root

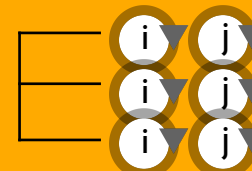


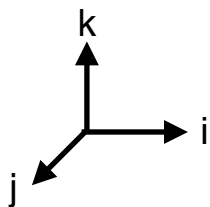
GPU  
Compatible  
Parallelization

physical core



physical core





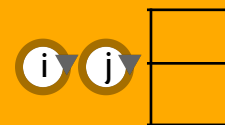
# Introduction | Hybrid Fortran

Hybrid Fortran  
approach

Hybrid  
Codebase

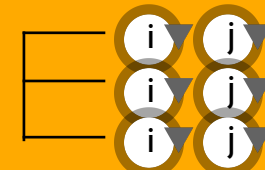
CPU  
Parallelization

physical core



GPU  
Parallelization

physical core



# Directives | Hybrid Fortran

```
1 module example
2 contains
3   subroutine wrapper(a, b, c, d)
4     real, intent(in) :: a(DOM(NX, NY, NZ)), b(DOM(NX, NY, NZ))
5     real, intent(out) :: c(DOM(NX, NY, NZ)), d(DOM(NX, NY, NZ))
6     integer(4) :: x, y
7
8     do y=1,NY
9       @parallelRegion{appliesTo(CPU), domName(x,y), domSize(NX, NY)}
10         call add(a(AT(x,y,:)), b(AT(x,y,:)), c(AT(x,y,:)))
11         call mult(a(AT(x,y,:)), b(AT(x,y,:)), d(AT(x,y,:)))
12       @end parallelRegion
13     end do
14   end subroutine
15
16   subroutine add(a, b, c)
17     real, intent(in) :: a(NZ), b(NZ)
18     real, intent(out) :: c(NZ)
19     integer :: z
20     @parallelRegion{appliesTo(GPU), domName(x,y), domSize(NX, NY)}
21     do z=1,NZ
22       c(z) = a(z) + b(z)
23     end do
24   @end parallelRegion
25 end module
```

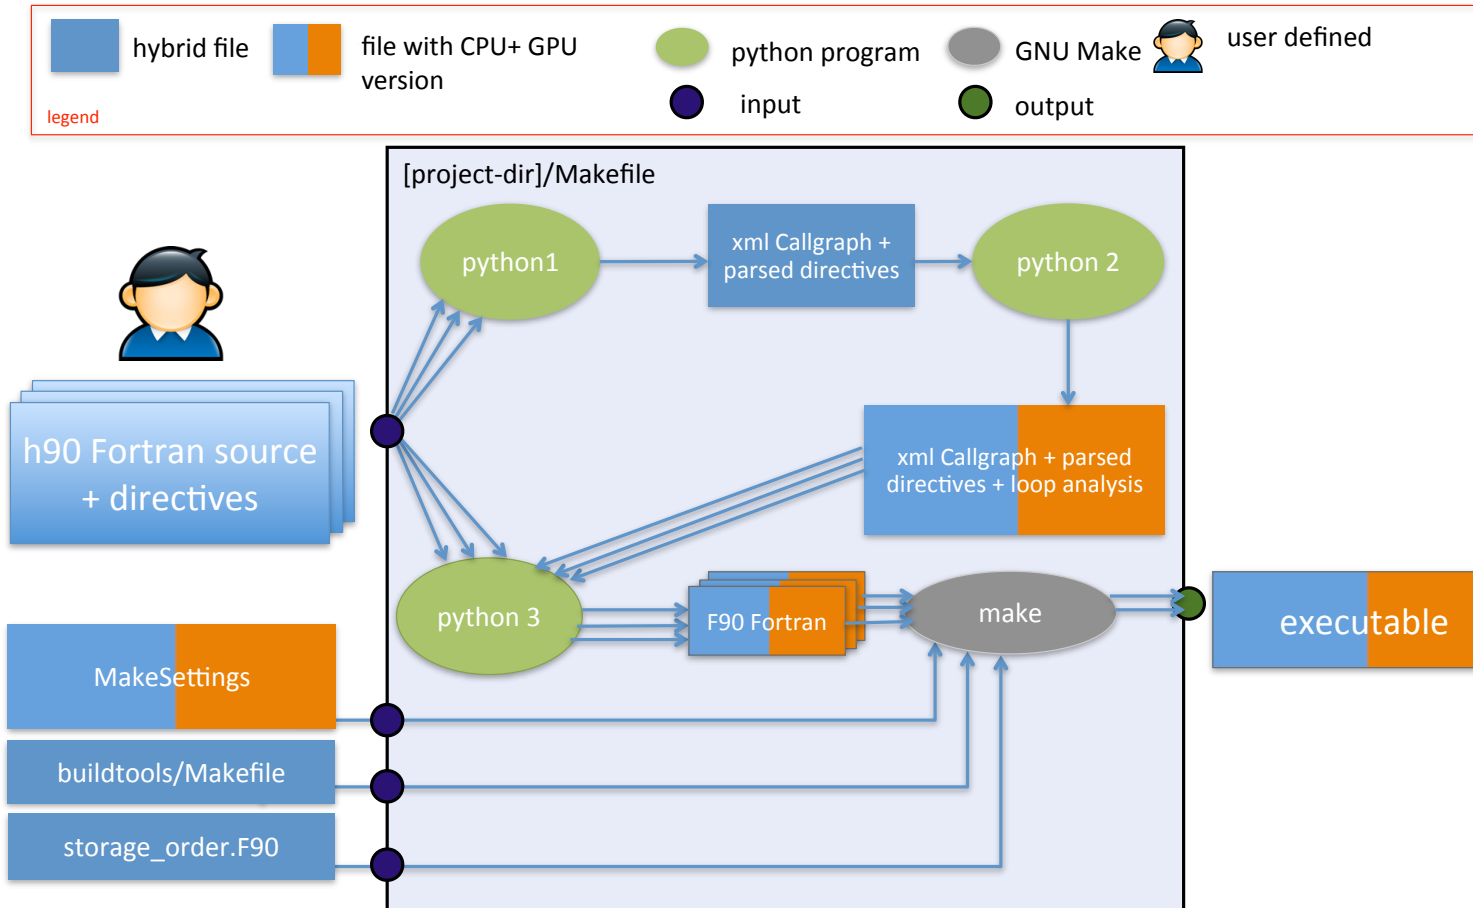
# Directives | Hybrid Fortran

```
module example
contains
  subroutine wrapper(a, b, c, d)
    real, dimension(NZ), intent(in) :: a, b
    real, dimension(NZ), intent(out) :: c, d
    real, dimension(NZ2), intent(out) :: e
    @domainDependant{domName(x,y), domSize(NX, NY), attribute(autoDom)}
    a, b, c, d, e
    @end domainDependant

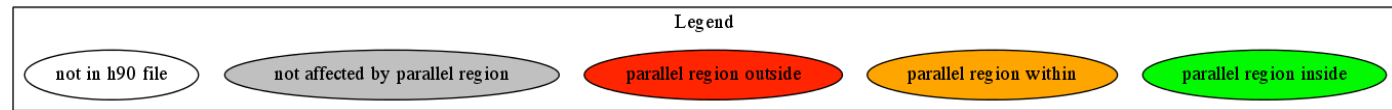
    @parallelRegion{appliesTo(CPU), domName(x, y), domSize(NX, NY)}
    call add(a, b, c)
    call mult(a, b, d)
    ...
    @end parallelRegion
  end subroutine
```



# Build System | Hybrid Fortran



# Scope of ASUCA PP Implementation



CPU  
Parallelization



Hybrid



GPU  
Parallelization

Motivation



Hybrid Fortran



ASUCA PP Implementation



Results & Discussion

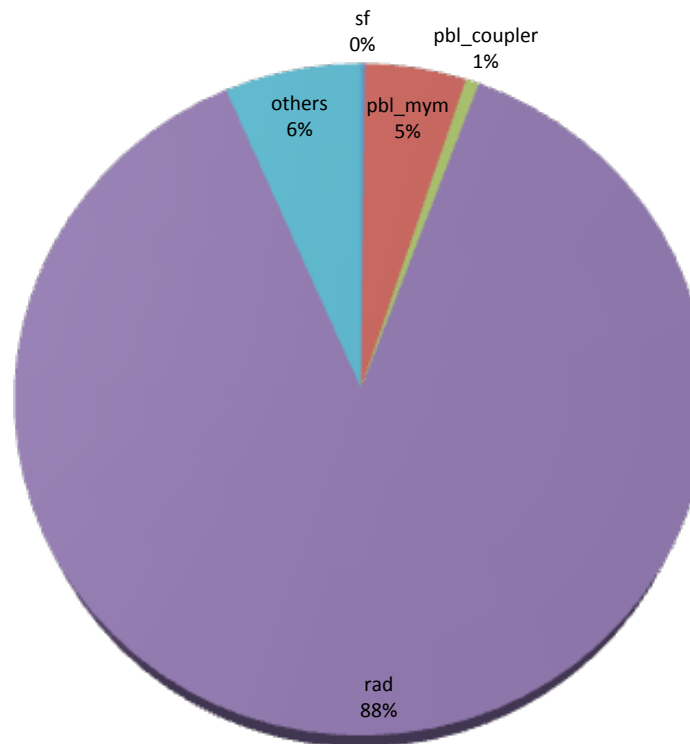


Conclusion

# Source Code Demo | Implementation

## Gabls3 Tests | Results & Discussion

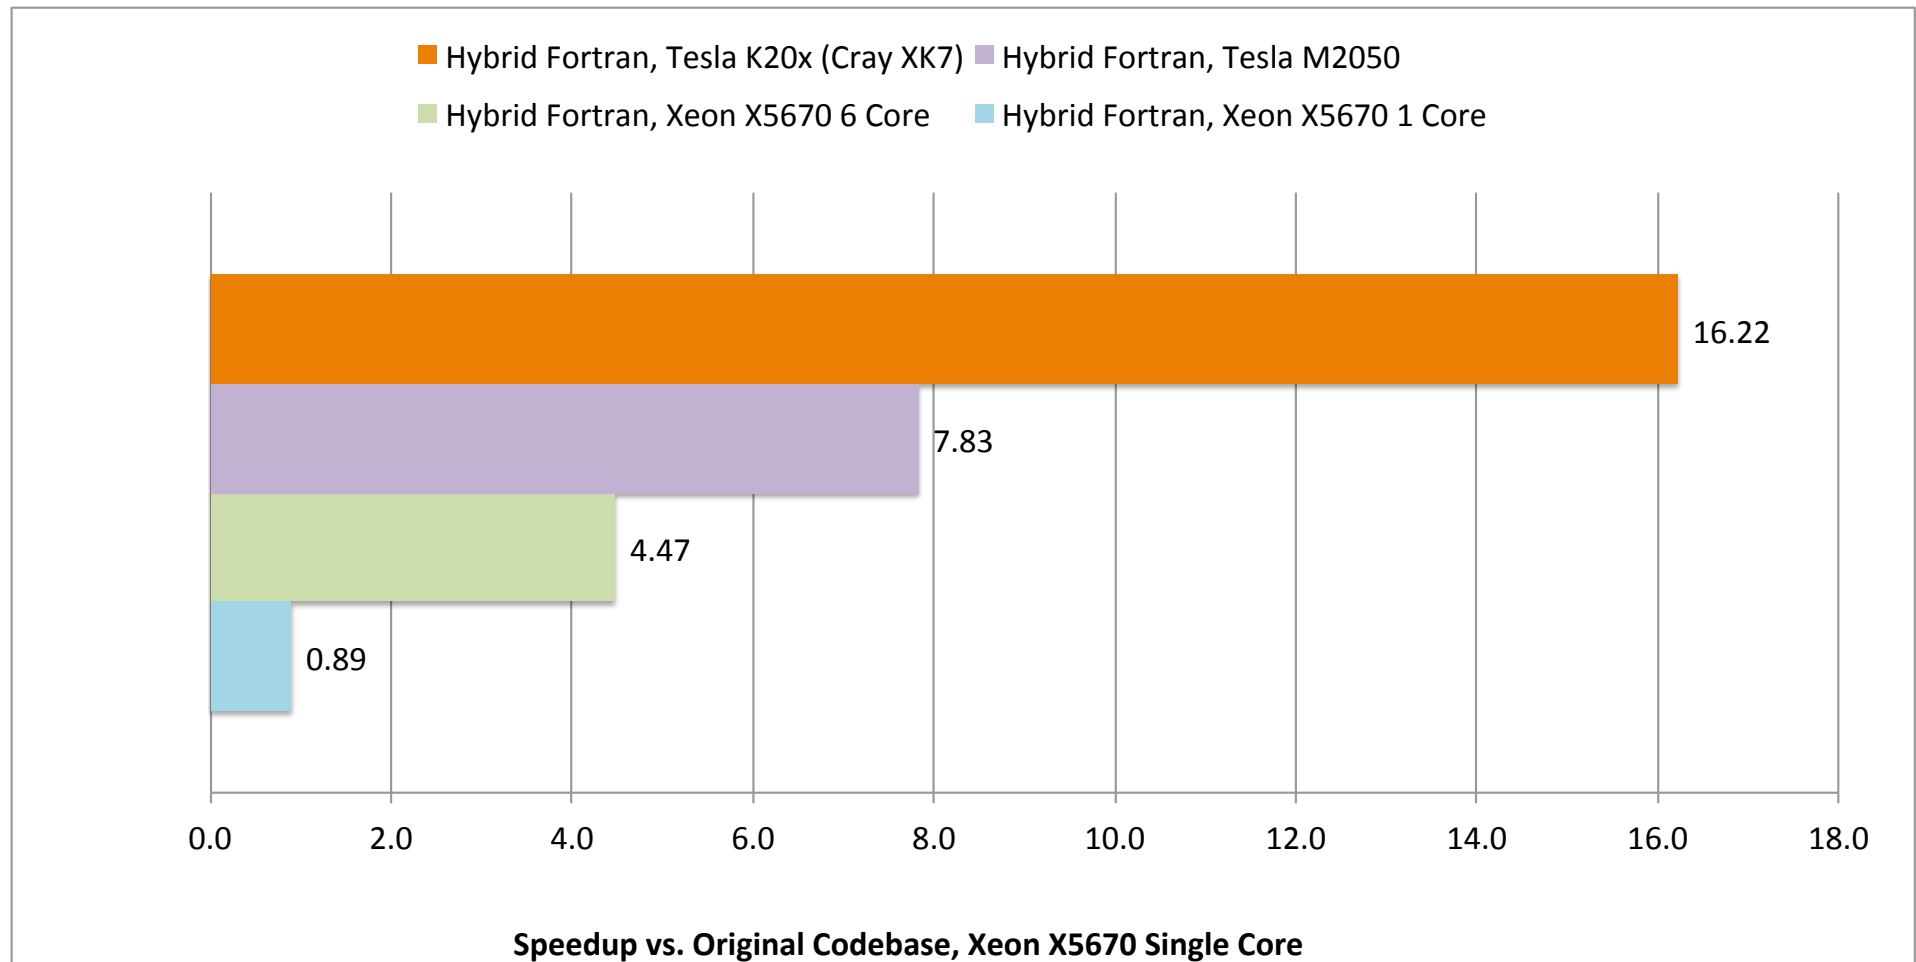
**Percentage of Runtime - Averaged over 100 timesteps on Westmere  
1 Core, Hybrid Fortran**



# Speedup Gabls3 vs. 1 Core

128 x 128 x 70, 100 Timesteps

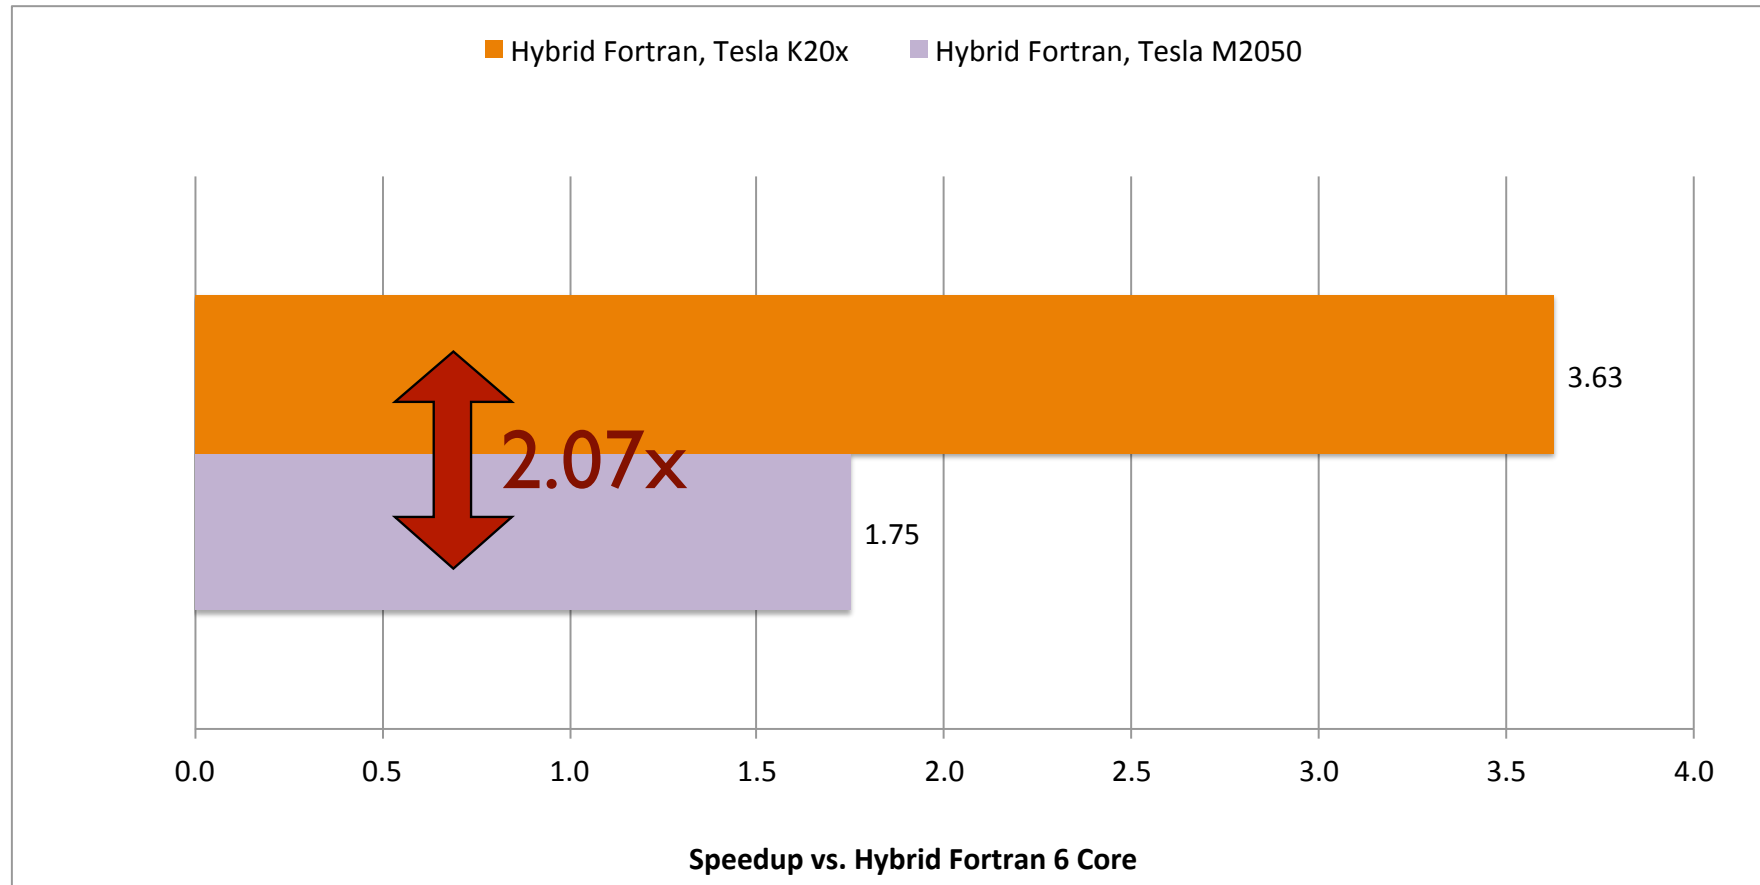
## Results & Discussion



# Speedup Gabls3 vs. 6 Core

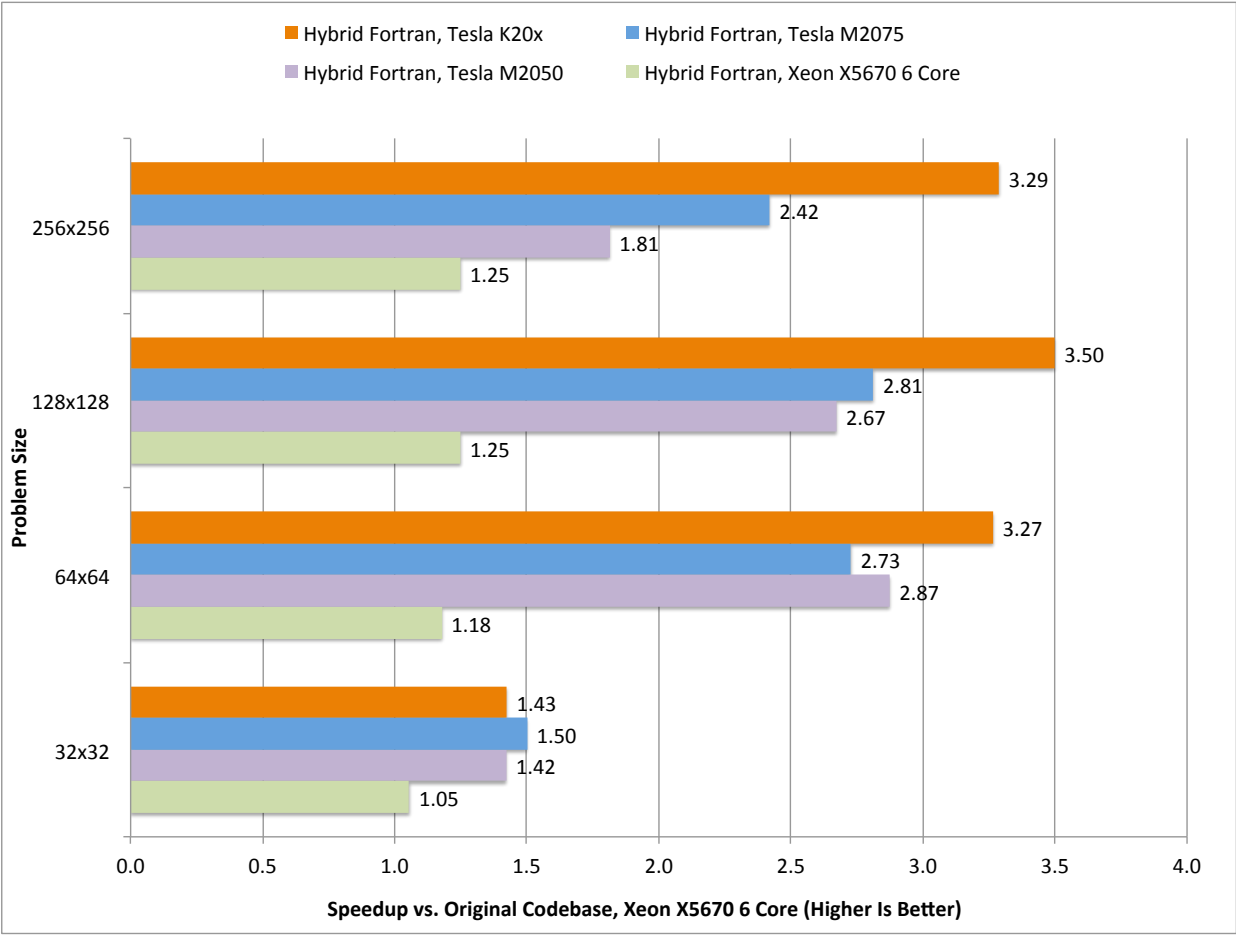
128 x 128 x 70, 100 Timesteps

## Results & Discussion



# Speedup Radiation vs. 6 Core , I x J x 63, Single Timestep

## Results & Discussion



Note: Speedup Single Core Hybrid Fortran vs. Single Core Original: 1-3%

# Why is HF Radiation Faster?

## Results & Discussion

- Original Implementation: Transmission Functions used in Longwave Radiation use ~4MB of Temporary Memory per Thread.
- This was too much for GPGPU, where  $\sim 10^5$  Threads are needed to saturate performance in this case.
- The Transmission therefore needed to be redesigned. Now only uses ~400KB per threads, a 10x improvement.



# Why is HF Radiation Faster?

Results & Discussion

Source Code

coefccs	1.776357E-15
coefc	1.081428E-14
cov	1.552159E-17
dlwbcs	6.82121E-13
dlwb	6.82121E-13
dswbcs	8.6402E-12
dswb	8.6402E-12
dswt	4.547474E-13
l_mo_inv	2.053913E-15
pt	1.136868E-13
qke	1.777826E-13
qsq	1.959569E-20
qv	7.314694E-17
qv_sfc	6.938894E-18
rh_sfc	0E+00
rlong	5.232738E-17
rnirb	0E+00
rnird	8.171241E-13
rshrt	1.12005E-16
rvisb	2.273737E-13
rvisd	<b>9.237056E-12</b>
tcvr	0E+00
tcwc	0E+00
tg	0E+00
tsq	2.097877E-14
ttranscs	3.885781E-16
ttrans	2.629026E-15
u	1.570067E-13
uf	7.494005E-16
ulwtcs	3.410605E-13
ulwt	3.410605E-13
uswbcs	4.902745E-13
uswb	4.902745E-13
uswtcs	6.593837E-12
uswt	6.593837E-12
v	9.878022E-15
wg	0E+00
zmean	4.440892E-16

# Accuracy

## Results & Discussion

after 450 Timesteps on GPU  
vs. JMA Reference

$$error = \sqrt{\sum_{1}^n (x_{HF} - x_r)^2}$$

# Programming Time

## Results & Discussion

~ 3M	Implementation of Hybrid Fortran
~ 3M	Implementation of ASUCA Physical Core

1. Hybridized ASUCA Physics with support for Single-GPU and CPU Multicore have been successfully completed.
2. Hybrid Fortran allows us to get GPU compatibility with baseline performance quickly, without compromising CPU performance.
3. We still have all the flexibility for performance tuning.

Questions

[michel@typhooncomputing.com](mailto:michel@typhooncomputing.com)